

Q. Build a simple linear regression model by performing EDA and do necessary transformations and select the best model using R or Python.

1) Delivery_time -> Predict delivery time using sorting time

1. Import Necessary Libraries

```
In [1]: import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")

import statsmodels.formula.api as smf
```

2. Import Data

```
In [6]: delivery_data = pd.read_csv("E:\Data Science by John\Assignments\Assignment 4- Simple Linear Regression\delivery_time.csv")
delivery_data
```

Out[6]:

	Delivery Time	Sorting Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10
5	15.35	6
6	19.00	7
7	9.50	3
8	17.90	10
9	18.75	9
10	19.83	8
11	10.75	4
12	16.68	7
13	11.50	3
14	12.03	3
15	14.88	4
16	13.75	6
17	18.11	7
18	8.00	2
19	17.83	7
20	21.50	5

3. Data Understanding

3.1 Perform Initial Analysis

In [7]: `delivery_data.shape`

Out[7]: (21, 2)

In [8]: `delivery_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Delivery Time    21 non-null     float64
1   Sorting Time     21 non-null     int64   
dtypes: float64(1), int64(1)
memory usage: 464.0 bytes
```

In [9]: `delivery_data.dtypes`

Out[9]: Delivery Time float64
Sorting Time int64
dtype: object

In [10]: `delivery_data.head()`

Out[10]:

	Delivery Time	Sorting Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10

```
In [11]: delivery_data.isna().sum()
```

```
Out[11]: Delivery Time    0  
Sorting Time    0  
dtype: int64
```

```
In [12]: delivery_data.describe()
```


```
Out[12]:
```

	Delivery Time	Sorting Time
count	21.000000	21.000000
mean	16.790952	6.190476
std	5.074901	2.542028
min	8.000000	2.000000
25%	13.500000	4.000000
50%	17.830000	6.000000
75%	19.750000	8.000000
max	29.000000	10.000000


3.2 Assumptions Check

1. Linearity

```
In [13]: sns.lmplot(x='Sorting Time', y='Delivery Time', data=delivery_data)
plt.xlabel('Sorting Time')
plt.ylabel('Delivery Time')
plt.show()
```



```
In [14]: sns.scatterplot(x='Sorting Time', y='Delivery Time', data=delivery_data)
plt.xlabel('Sorting Time')
plt.ylabel('Delivery Time')
plt.show()
```



```
In [15]: delivery_data.skew()
```

```
Out[15]: Delivery Time    0.352390
Sorting Time    0.047115
dtype: float64
```


```
In [16]: delivery_data.kurt()
```

```
Out[16]: Delivery Time    0.317960
Sorting Time   -1.148455
dtype: float64
```

Linearity test is Failed

2. Normality

```
In [17]: sns.distplot(a= delivery_data['Sorting Time'],hist = False)  
plt.show()
```



Normality test is Failed

3. No Multicollinearity II 4. No Autoregression

These test are passed because we have only one input so No Multicollinearity & We don't have datetime datatype of our data

4. Data Preparation

```
In [18]: delivery_data.columns = [column.replace(" ", "_") for column in delivery_data.columns]
delivery_data.head()
```

```
Out[18]:
```

	Delivery_Time	Sorting_Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10

5. Model Building II 6. Model Training

Linear Regression model can be built by using any of the 2 libraries:

1. Statmodels
2. sklearn

Using Statsmodels:

```
In [19]: linear_model_sm = smf.ols(formula = 'Delivery_Time ~ Sorting_Time', data = delivery_data).fit()
linear_model_sm
```

```
Out[19]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x259a0d6ba60>
```

```
In [20]: linear_model_sm.params
```

```
Out[20]: Intercept      6.582734
Sorting_Time    1.649020
dtype: float64
```


7. Model Testing

```
In [21]: delivery_data.head()
```

```
Out[21]:
```

	Delivery_Time	Sorting_Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10

```
In [22]: Delivery_Time_Pred = linear_model_sm.predict(delivery_data['Sorting_Time'])  
Delivery_Time_Pred.head()
```

```
Out[22]: 0    23.072933  
1    13.178814  
2    16.476853  
3    21.423913  
4    23.072933  
dtype: float64
```

8. Model Evaluation


```
In [23]: error = delivery_data['Delivery_Time'] - Delivery_Time_Pred
error
```

```
Out[23]: 0    -2.072933
         1     0.321186
         2     3.273147
         3     2.576087
         4     5.927067
         5    -1.126853
         6     0.874127
         7    -2.029794
         8    -5.172933
         9    -2.673913
        10     0.055107
        11    -2.428814
        12    -1.445873
        13    -0.029794
        14     0.500206
        15     1.701186
        16    -2.726853
        17    -0.015873
        18    -1.880774
        19    -0.295873
        20     6.672167
dtype: float64
```

.....back to Assumption Check

5. Homoscedasticity Check


```
In [24]: plt.scatter(x= delivery_data['Sorting_Time'], y=error)
plt.title('Sorting_Time Vs Error')
plt.xlabel('Sorting_Time')
plt.ylabel('Error')
plt.show()
```



Homoscedasticity Check is Passed

6. Zero Residual Mean across the Fitted Line.

```
In [25]: plt.scatter(x=delivery_data['Delivery_Time'],y=Delivery_Time_Pred)
plt.title('Actual O/P Vs Predicted O/p')
plt.xlabel('Actual O/P')
plt.ylabel('Predicted O/p')
plt.show()
```



Zero Residual Mean Test is Failed

8.1 Evaluation Metrics of Linear Regression

```
In [26]: print('R2Score      : ',linear_model_sm.rsquared.round(4)) #Overall Contribution of Predictors
print('Adj.R2Score   : ',linear_model_sm.rsquared_adj.round(4)) #Overall Contribution of Predictors
print('AIC Value    : ',linear_model_sm.aic.round(4)) #Error Impurity
print('BIC Value    : ',linear_model_sm.bic.round(4)) #Error Impurity
```

```
R2Score      : 0.6823
Adj.R2Score   : 0.6655
AIC Value     : 106.714
BIC Value     : 108.803
```

9. Model Prediction

9.1 Manual Prediction

```
In [28]: linear_model_sm.params
```

```
Out[28]: Intercept      6.582734
Sorting_Time    1.649020
dtype: float64
```

```
In [27]: # Manual prediction for sorting time 10
delivery_time = (6.582734) + (1.649020)*(10)
delivery_time
```

```
Out[27]: 23.072933999999997
```

```
In [29]: # Manual prediction for sorting time 20
delivery_time = (6.582734) + (1.649020)*(20)
delivery_time
```

```
Out[29]: 39.563134
```

9.2 Automatic Prediction using model

```
In [36]: delivery_data.head()
```

```
Out[36]:
```

	Delivery_Time	Sorting_Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10

```
In [37]: test_data = pd.DataFrame(data = {"Sorting_Time":[10,20]})
test_data
```

```
Out[37]:
```

	Sorting_Time
0	10
1	20

```
In [38]: linear_model_sm.predict(test_data)
```

```
Out[38]: 0    23.072933
1    39.563132
dtype: float64
```