

Q. Build a simple linear regression model by performing EDA and do necessary transformations and select the best model using R or Python.

2) Salary_hike -> Build a prediction model for Salary_hike

1. Import Necessary Libraries ¶

```
In [38]: import pandas as pd

import seaborn as sns
from matplotlib import pyplot as plt

import warnings
warnings.filterwarnings("ignore")

from sklearn.linear_model import LinearRegression
import statsmodels.formula.api as smf

from sklearn.metrics import mean_squared_error, mean_absolute_error
```

2. Import Data

```
In [2]: salary_data = pd.read_csv("E:\Data Science by John\Assignments\Assignment 4- Simple Linear Regression\Salary_Data.csv")
salary_data
```

Out[2]:

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0

	YearsExperience	Salary
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

3. Data Understanding

3.1 Perform Initial Analysis

In [3]: salary_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```

In [4]: salary_data.shape

Out[4]: (30, 2)

```
In [5]: salary_data.describe()
```

```
Out[5]:
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000


```
In [6]: salary_data.isna().sum()
```

```
Out[6]: YearsExperience    0  
Salary                    0  
dtype: int64
```

3.2 Assumptions Check

1. Linearity Check


```
In [7]: sns.lmplot(x="YearsExperience", y="Salary", data=salary_data)
plt.xlabel("YearsExperience")
plt.ylabel("Salary")
plt.show()
```



Linearity Test is passed

2. Normality Check

```
In [8]: sns.distplot(a= salary_data['YearsExperience'], hist = False)  
plt.show()
```



Normality Test is Failed

3. Multicollinearity || 4. Autoregression

These test are passed because we have only one input so No Multicollinearity & We don't have datetime datatype of our data

4. Data Preparation

Already in required format so no need to any changes

5. Model Building || 6. Model Training

Using Sklearn

```
In [9]: x = salary_data[['YearsExperience']]  
y = salary_data[['Salary']]
```

```
In [10]: linear_model = LinearRegression()
```

```
In [11]: linear_model_sk = linear_model.fit(x,y)
```

```
In [12]: linear_model.coef_
```

```
Out[12]: array([[9449.96232146]])
```

```
In [13]: linear_model.intercept_
```

```
Out[13]: array([25792.20019867])
```

7. Model Testing

Manual Prediction

```
In [14]: #  $y = mx + c$   
# If  $x = 2$ ,  $y = ??$   
(9449.96232146 * 2) + 25792.20019867
```

```
Out[14]: 44692.124841590005
```

```
In [15]: #  $y = mx + c$   
# If  $x = 5$ ,  $y = ??$   
(9449.96232146 * 5) + 25792.20019867
```

```
Out[15]: 73042.01180597
```

```
In [16]: # y = mx+c  
# If x = 6.5, y = ??  
(9449.96232146 * 6.5) + 25792.20019867
```

Out[16]: 87216.95528816

```
In [17]: # y = mx+c  
# If x = 8, y = ??  
(9449.96232146 * 8) + 25792.20019867
```

Out[17]: 101391.89877035

```
In [18]: # y = mx+c  
# If x = 10, y = ??  
(9449.96232146 * 10) + 25792.20019867
```

Out[18]: 120291.82341327

Automatic Prediction

```
In [19]: y.head()
```

Out[19]:

	Salary
0	39343.0
1	46205.0
2	37731.0
3	43525.0
4	39891.0


```
In [20]: X_test = pd.DataFrame(data={'YearsExperience': [2, 5, 6.5, 8, 10]})  
X_test
```

```
Out[20]:
```

	YearsExperience
0	2.0
1	5.0
2	6.5
3	8.0
4	10.0

```
In [21]: linear_model_sk.predict(X_test)
```

```
Out[21]: array([[ 44692.12484158],  
                [ 73042.01180594],  
                [ 87216.95528813],  
                [101391.89877031],  
                [120291.82341322]])
```

```
In [22]: y_pred = linear_model_sk.predict(x)
y_pred
```

```
Out[22]: array([[ 36187.15875227],
 [ 38077.15121656],
 [ 39967.14368085],
 [ 44692.12484158],
 [ 46582.11730587],
 [ 53197.09093089],
 [ 54142.08716303],
 [ 56032.07962732],
 [ 56032.07962732],
 [ 60757.06078805],
 [ 62647.05325234],
 [ 63592.04948449],
 [ 63592.04948449],
 [ 64537.04571663],
 [ 68317.03064522],
 [ 72097.0155738 ],
 [ 73987.00803809],
 [ 75877.00050238],
 [ 81546.97789525],
 [ 82491.9741274 ],
 [ 90051.94398456],
 [ 92886.932681  ],
 [100446.90253816],
 [103281.8912346 ],
 [108006.87239533],
 [110841.86109176],
 [115566.84225249],
 [116511.83848464],
 [123126.81210966],
 [125016.80457395]])
```

8. Model Evaluation

```
In [23]: error = y - y_pred  
error
```

Out[23]:

	Salary
0	3155.841248
1	8127.848783
2	-2236.143681
3	-1167.124842
4	-6691.117306
5	3444.909069
6	6007.912837
7	-1587.079627
8	8412.920373
9	-3568.060788
10	570.946748
11	-7798.049484
12	-6635.049484
13	-7456.045717
14	-7206.030645
15	-4159.015574
16	-7958.008038
17	7210.999498
18	-183.977895
19	11448.025873
20	1686.056015
21	5386.067319
22	855.097462
23	10530.108765

	Salary
24	1424.127605
25	-5259.861092
26	1402.157748
27	-3876.838485
28	-735.812110
29	-3144.804574

.....back to Assumption Check

5. Homoscedasticity Check

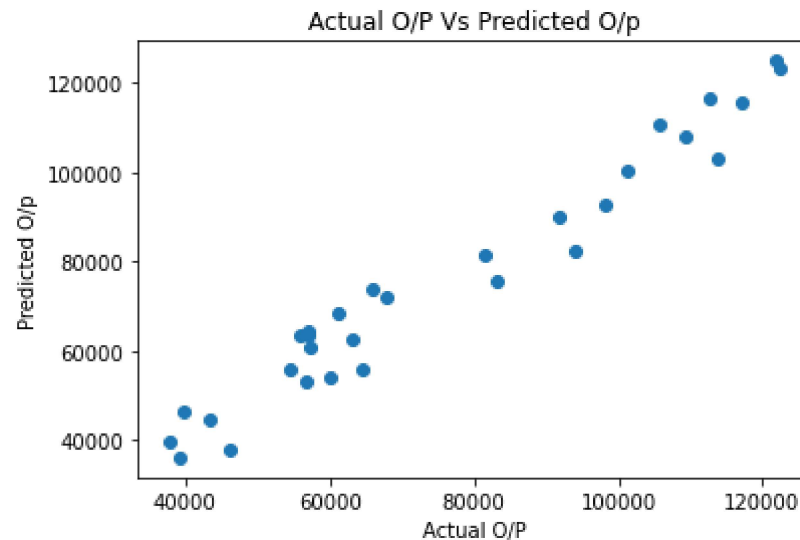
```
In [24]: plt.scatter(x= salary_data['YearsExperience'], y=error)
plt.title('YearsExperience Vs Error')
plt.xlabel('YearsExperience')
plt.ylabel('Error')
plt.show()
```



Homoscedasticity Check is Passed

6. Zero Residual Mean across the Fitted Line.

```
In [25]: plt.scatter(x=salary_data['Salary'],y=y_pred)
plt.title('Actual O/P Vs Predicted O/p')
plt.xlabel('Actual O/P')
plt.ylabel('Predicted O/p')
plt.show()
```



Zero Residual Mean Test is Passed

Build Linear Model using StatsModels

Model Building || Model Training

```
In [27]: salary_data.head()
```

```
Out[27]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
In [29]: linear_model_stats = smf.ols(formula = 'Salary~YearsExperience', data = salary_data).fit()  
linear_model_stats
```

```
Out[29]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x20bed656ac0>
```

```
In [30]: linear_model_stats.params
```

```
Out[30]: Intercept          25792.200199  
YearsExperience      9449.962321  
dtype: float64
```

8.1 Evaluation Metrics of Linear Regression

```
In [31]: print('R2Score      : ',linear_model_stats.rsquared.round(4)) #Overall Contribution of Predictors
print('Adj.R2Score  : ',linear_model_stats.rsquared_adj.round(4)) #Overall Contribution of Predictors
print('AIC Value    : ',linear_model_stats.aic.round(4)) #Error Impurity
print('BIC Value    : ',linear_model_stats.bic.round(4)) #Error Impurity
print('P-Value      \n:',linear_model_stats.pvalues)
```

```
R2Score      : 0.957
Adj.R2Score  : 0.9554
AIC Value    : 606.8823
BIC Value    : 609.6847
P-Value
: Intercept      5.511950e-12
YearsExperience  1.143068e-20
dtype: float64
```

```
In [39]: mean_squared_error(y,y_pred)
```

```
Out[39]: 31270951.722280968
```

```
In [40]: mean_absolute_error(y,y_pred)
```

```
Out[40]: 4644.2012894435375
```



```
In [33]: y_pred_stats = linear_model_stats.predict(x)
y_pred_stats
```

```
Out[33]: 0      36187.158752
1      38077.151217
2      39967.143681
3      44692.124842
4      46582.117306
5      53197.090931
6      54142.087163
7      56032.079627
8      56032.079627
9      60757.060788
10     62647.053252
11     63592.049484
12     63592.049484
13     64537.045717
14     68317.030645
15     72097.015574
16     73987.008038
17     75877.000502
18     81546.977895
19     82491.974127
20     90051.943985
21     92886.932681
22    100446.902538
23    103281.891235
24    108006.872395
25    110841.861092
26    115566.842252
27    116511.838485
28    123126.812110
29    125016.804574
dtype: float64
```

```
In [42]: error_stats = salary_data['Salary'] - y_pred_stats  
error_stats
```

```
Out[42]: 0      3155.841248  
1      8127.848783  
2     -2236.143681  
3     -1167.124842  
4     -6691.117306  
5      3444.909069  
6      6007.912837  
7     -1587.079627  
8      8412.920373  
9     -3568.060788  
10      570.946748  
11     -7798.049484  
12     -6635.049484  
13     -7456.045717  
14     -7206.030645  
15     -4159.015574  
16     -7958.008038  
17      7210.999498  
18     -183.977895  
19     11448.025873  
20      1686.056015  
21      5386.067319  
22       855.097462  
23     10530.108765  
24      1424.127605  
25     -5259.861092  
26      1402.157748  
27     -3876.838485  
28     -735.812110  
29     -3144.804574  
dtype: float64
```

